



AGENTIC SYSTEMS

**Your agents
believe
each other.**

#2 in the silent-failure series



THE STAKES

One bad value, and everything downstream inherits it.

Not as an error. As a fact,
indistinguishable from the ones
that were derived honestly.




WHAT WE ALL ASSUME



Validate at the edge and you are covered.

Check the input, sanitise the prompt, filter what the tool returns. This is what a guardrail means to most teams.




THE BREAK

A perimeter needs exactly one way in.

In a fleet passing structured
output between agents, every
hop is a way in — so you would
need one perimeter per edge.


and nobody draws that graph



WHERE IT STOPS BEING CHECKED



Checked once, at the front.
Every arrow after that is a
handover nothing re-inspects.



SOLVED ONCE ALREADY

1976 Denning

**Information flow
between classes
forms a lattice.**

And a program's flows can be
certified against it
mechanically, rather than
reviewed by hand.

A lattice model of secure information flow · Comm. ACM
19(5)



AND AGAIN


1997 Myers & Liskov

Mutual distrust, decentralised authority.

Hand data to code you do not
trust, and still govern what it
may do with it downstream.

Same shape as a fleet.

they were describing programs, not agents · the shape
transfers, the claim does not




WHAT BIOLOGY DOES

The body does not defend itself at the skin.

Every cell continuously
presents what it is carrying.
Non-self is caught inside, at
every hop, rather than at one
wall.

a parallel, not a theorem · stated as a shape



THE TEST

How many hops before anything re-checks a value?

If an agent cannot say where a value came from, it cannot answer that — and neither can you.

run it against your own fleet



WHERE THIS STOPS WORKING

Provenance costs something on every hop.

On a two-agent pipeline where both sides are yours, a perimeter is the right size of answer. The cost lands when the graph grows.

the rule has an edge, and this is it