




AUTONOMY LEVEL 5

**Each autonomy
level is defined
by the evidence
it requires.**


the level is set by what can be shown to somebody who does not trust the operator



THE DISTINCTION

**Level 2 and level 4
can look identical
from outside.**

The behaviour is the same. What differs is what can be demonstrated to somebody who does not trust the operator.



L1

**Suggests. A
human performs
every action.**

The evidence is trivial because
nothing acted.

Acts reversibly.

Every effect can be undone. The evidence is the undo path, together with a record that the undo was exercised rather than assumed.



L3

Acts under a gate.

A human approves the irreversible operations. The evidence is that the gate cannot be circumvented, including by the system's own retry logic.


this is where it usually fails without anyone noticing

Acts, audited afterwards.

The evidence is a record complete enough that an auditor can reconstruct any single decision after the fact, which a log by itself does not provide.

Acts unattended.

The evidence is an enumerated set of failure modes that are not covered, with a stated reason for accepting each one.



THE LADDER

Every rung above the one you can prove is decoration.

L5 · unattended

L4 · audited after

L3 · acts, gated

L2 · acts, reversibly

L1 · suggests

dashed rungs are the honest rendering of evidence that does not exist




WHERE THE COST SITS

The cost sits in the proving.

Building a system that acts
without approval takes days.

Producing the record that lets an
auditor reconstruct why it acted
takes considerably longer, and it
is the part that gets deferred.



THE TEST

**Which level would
your system survive
an external audit at?**

The level it could be defended at
by somebody assuming the
operator is wrong, which is
commonly one rung below the
level it runs at.

run it against something already in production