

- Agent systems

Memory that can change.

What should an intelligent product remember when its user changes?

Said	Keep this comparison short.
Inferred	Prefers short answers.
Scope	Comparisons on mobile

- Agent systems

The design question

A useful memory gives someone a place to continue. It also gives them a way to become a different user. Put these statements beside each other: “Keep this comparison short” and “Prefers short answers.” The first records a request. The second turns that request into a general interpretation. They can produce very different experiences when the person returns for a detailed tutorial.

- Agent systems

The design question


That gap is where I would begin designing an intelligent product's memory. The interesting problem is deciding what should influence the next interaction, and making that decision correctable.

- Agent systems

A record is not yet a relationship

Start with a small working object: a memory card that shows the original statement, the current interpretation and the context in which it applies. In this example, the person could change the scope to “comparisons on mobile.” The product can then keep those comparisons concise while allowing a tutorial to develop an explanation.

- Agent systems



A record is not yet a relationship

The correction is useful because it changes what the software does next. A memory-settings page that merely displays the stored sentence would leave the central interpretation untouched. The person needs a way to reshape the assumption, not just admire the archive.

- Agent systems

A record is not yet a relationship

Said	Keep this comparison short.
Inferred	Prefers short answers.
Scope	Comparisons on mobile

Figure 1. An illustrative memory card. The scope determines where an interpretation may be useful; these are proposed interface fields, not measured user outcomes.

- Agent systems

Selection has a purpose

A historical research example helps make the engineering choice visible. The 2023 Generative Agents architecture used relevance, recency and importance to retrieve records, and reflection to form higher-level inferences. Its evaluation concerned believable behavior in a simulated world. It did not establish a universal recipe for remembering real people. [Park et al., 2023.](#)

- Agent systems



Selection has a purpose

The transferable question is why a record should be brought back now. A planning product might need the unfinished decision and the constraint behind it. A learning workspace might need the project that still interests the learner. A complete miniature of the previous conversation is not necessarily the right working object for either.

- Agent systems

Selection has a purpose

The paper's retrieval idea can be written as a weighted score: $\text{score}(m) = w_R \cdot R(m) + w_C \cdot C(m) + w_I \cdot I(m)$. Here m is a memory; R , C and I are normalized relevance, recency and importance scores; the w terms are weights. The score is dimensionless. Raising a weight changes which consideration has more influence, but does not make the selected record true or appropriate. A product must choose and investigate its own retrieval policy.

- Agent systems

Selection has a purpose



Figure 2. Selection connects stored records to a particular task. This conceptual diagram makes no claim about retrieval accuracy or model performance.

- Agent systems

Preserve unfinished thinking


Consider a second kind of record: “We chose a guided first project because the learner wanted to begin without choosing a direction.” Add a condition for reopening it: “They bring a project idea of their own.” Keep the open-workspace alternative nearby.

- Agent systems

Preserve unfinished thinking

Now the product has something more useful than a preference label. It has a decision, a reason and an occasion to reconsider. If the learner returns with an idea, the software can offer a different entrance instead of insisting on the route that previously helped.

- Agent systems



Preserve unfinished thinking

This example also exposes an important distinction. The old memory may remain accurate while losing its usefulness. “That is wrong” and “stop using that here” should therefore be separate ways to revise the relationship.

- Agent systems



Prototype the disagreement

I would test a memory interface by asking someone to continue their own unfinished work, then change a meaningful constraint. Watch whether they can express the change, inspect the interpretation it creates and keep the parts that still matter. Include starting fresh as a real option.

- Agent systems

Prototype the disagreement

A task that happens once may not benefit from persistent memory at all. A long-running project may need substantial continuity. The design should follow that difference rather than treating accumulation as inherently valuable.

- Agent systems

Prototype the disagreement

There is also a cost to making every assumption editable: the person can end up administering a database of themselves. Start with the assumptions that materially change the experience, and put revision beside the affected work. Do not demand a classification exercise before the person can continue.

- Agent systems

Prototype the disagreement

The promise I would make is concrete: the product helps you resume, shows what it is carrying forward and lets you change its mind about you. That is a relationship worth prototyping.